


**We saved \$400/month
on day one by NOT
using the obvious choice.**

The full LLMWiki architecture on AWS
and every decision behind it.



Swipe → full architecture breakdown, every decision explained

THE PROBLEM

This is how most enterprises manage critical knowledge today.

- Critical docs buried in shared drives nobody searches
- Tribal memory — one person holds the answer
- Email thread to find the expert — 2 days wait
- Answer walks out the door when they resign


Sound familiar? Swipe to see the fix →

THE INGESTION PIPELINE

Any format. Any size.
3 minutes to queryable.



Amazon S3 ← EventBridge trigger on upload



PDF → Markdown

Chunk + Embed

Vector Store ■ ~\$0/mo base

✓ Pure event-driven. Zero idle cost.

No polling. No scheduled jobs. Fires on every upload.

THE QUERY PIPELINE

Question in. Cited answer out.



Confidence Score attached to every answer:

HIGH

2+ direct-match sources

MEDIUM

Partial match sources

LOW

Gap Detection triggered

Every answer includes the source wiki pages that were used.

No hallucination without citation. Gaps are flagged, not fabricated.

THE DATA LAYER

Every query logged.
Every source tracked.

llmwiki-index

Page registry — every wiki document, · its slug, type, tags, confidence baseline

llmwiki-log

Full audit trail — every query, every · agent invocation, every confidence score

llmwiki-source-registry

Document provenance — uploader, timestamp, · trust level, KB sync status

Compliance-ready audit trail out of the box.

THE COST DECISION

Same capability.
 $1/\infty$ the floor cost.

OpenSearch**\$700+/month**

minimum — whether
you use it or not

Always-on capacity
OCU billing
Complex IAM setup

VS

■ S3 Vectors**~\$0/month base**

pay per query
zero idle cost

Serverless native
S3 IAM — already set up
Same vector search

We chose S3 Vectors on day one. Still the right call.

THE FIVE PILLARS

Production, not a demo.

■ DISTRIBUTED

Event-driven ingestion + retrieval
Independent components — no bottleneck

■ FAULT-TOLERANT Idempotent Lambdas · DLQs · DynamoDB state
Automatic retry — no manual intervention

■ SECURED

KMS encryption at rest + in transit
IAM least-privilege per Lambda · VPC isolation

■ SCALABLE

Serverless-first — scales to zero
Scales to millions of documents · no rearchitecting

■ VALIDATED

22 unit tests generated by Codex
First-pass: 22 passed · 0 failed · CI-ready

WHAT'S NEXT

This architecture just got a brain.

Next week: Neuro-SAN & the AAOSA Protocol

5 specialist agents collaborating on one customer question.

No agent ever sees the sensitive data.

Orchestrated in HOCON — no Python business logic required.

Want to see this on your own documents?

30 minutes · Live demo · DM me on LinkedIn